

Conscious Consistency with Akka Cluster, CRDTs and Distributed Data.

Jan Pustelnik  @gosubpl

6 April 2017

```
DIV10:  
SBC #10 ;  
Y=INT(X/10)  
INY  
BCS DIV10  
ADC #10+'0'  
CPY #"0"  
BEQ ONEDIG  
PHA ; IF Y>0 THEN  
TYA ? Y;  
JSR PUTCH  
PLA ; ? X MOD 10;
```



akka / akka

<> Code



gosubpl

24 commits / 7,081 ++ / 1,060 --

#31

ORMap and friends have deltas (#22350)

master (#22508) v2.5.0-RC2 v2.5.0-RC1



gosubpl committed on 23 Feb

Showing 15 changed files with 4,520 additions and 289 deletions.

■ CRDTs

■ Why

■ What / How

■ in akka

Why?

GFT



- **Why?**

Why?

GFT



■ **CRDTs**

- CRDTs

- „C”

- CRDTs

- „C”

- ???

- **CRDTs**

- „C”

- ???

- **In a minute...**

- **CRDTs**
- **Replicated**

- **CRDTs**
 - **Replicated**
- **Cluster, Replication**


- **CRDTs**
 - **Replicated**
- **Cluster, Replication**
 - **Cluster provides very low level primitives – nodes, communication primitives**

- **CRDTs**
 - **Replicated**
- **Cluster, Replication**
 - **Cluster provides very low level primitives – nodes, communication primitives**
 - **There is a Replicator in the Cluster**

- **CRDTs**
- **Data Types**

- **CRDTs**
 - **Data Types**
- **„Data Types” most obvious here 😊**

Why?



Cannot be fully available, global consistency (CP)

HAT – session based consistency, some availability limitations may apply

Cannot be fully consistent (AP)

- **CRDTs**
 - **Concurrent**
 - **???**

- **CRDTs**
- **Concurrent**


- **CRDTs**

- **Concurrent**

- **Allow concurrent updates on Cluster nodes without requiring synchronisation (AP)**

- **CRDTs**
 - **Concurrent**
 - **High Availability**
 - **High Throughput**
 - **Great for Shopping Cart!**


Why?




Good Ol' SQL –
„ACID” 😊
CRDTs - NO!

CRDTs ?

NoSQL! CRDTs ?



Anomalies! We don't want to be here!




(a) Dynamo shopping cart

An optimized conflict-free replicated set

Bieniusa A., et al.


<https://arxiv.org/abs/1210.3368v1>



(a) Dynamo shopping cart

Shopping cart service: 99.94% of requests saw 1 version


<http://s3.amazonaws.com/AllThingsDistributed/sosp/amazon-dynamo-sosp2007.pdf>



(a) Dynamo shopping cart

Shopping cart service: 99.94% of requests saw 1 version

0.06% saw 2 versions...




(a) Dynamo shopping cart

Shopping cart service: 99.94% of requests saw 1 version

That makes 6 out of 10 thousand cases where we ship a book that hasn't been paid for...

Why?



CRDTs!

Anomalies! We don't want to be here!

- **CRDTs**
 - **Convergent**
 - **Conflict-free**

- **CRDTs**
 - **Convergent**
 - **Any two replicas will converge to the same state after updates have stopped**

■ CRDTs


■ Conflict-free

- **State change is incremental, no historical state is changed – a (potential) conflict once resolved stays resolved**

■ What/How?

- **GSet – Grow-only Set**
 - **API – only 2 operations:**
 - **add(element)**
 - **elements()**
 - **No removals!**

GSet – Timeline



Based on *A comprehensive study of Convergent and Commutative Replicated Data Types* – Shapiro et al.
<https://hal.inria.fr/inria-00555588/document>

- **GSet – Grow-only Set**
 - **API – only 2 operations:**
 - **add(element)**
 - **elements()**
 - **No removals!**
 - **Update is merged with existing set with standard Set Union operation**

- **GSet – Grow-only Set**
 - **Update is merged with existing set with standard Set Union operation**

- **GSet – Grow-only Set**
 - **Update is merged with existing set with standard Set Union operation**
 - **Wait a second, what's an update?**

- **GSet – Grow-only Set**
 - **Update is merged with existing set with standard Set Union operation**
 - **Wait a second, what's an update?**
 - **Update contains the full state of GSet (all elements!)**

- **GSet – Grow-only Set**
 - **Update is merged with existing set with standard Set Union operation**
 - **Wait a second, what's a merge?**

- **GSet – Grow-only Set**
 - **Update is merged with existing set with standard Set Union operation**
 - **Wait a second, what's a merge?**
 - **Merge is Set Union, which is commutative, idempotent and associative.**

- **GSet – Grow-only Set**
 - **This are the consequences of GSet being a monotonic join semilattice object**
 - **Which means that its state cannot decrease and merge results in calculating Least Upper Bound of source states**




- **In our case of GSet the Least Upper Bound is the smallest set containing all updates, i.e. it is Set Union of all updates 😊**
- **Least Upper Bound operation can generally be shown to be commutative, idempotent and associative**

- **GCounter – Grow-only Counter**
 - **Two operations:**
 - **Increment(x)**
 - **GetValue**
 - **Merge requires identification of node and doing max over all nodes**

Specification 6 State-based increment-only counter (vector version)

- 1: payload integer $[n]$ P ▷ One entry per replica
 - 2: initial $[0, 0, \dots, 0]$
 - 3: update *increment* ()
 - 4: let $g = myID()$ ▷ g : source replica
 - 5: $P[g] := P[g] + 1$
 - 6: query *value* () : integer v
 - 7: let $v = \sum_i P[i]$
 - 8: compare (X, Y) : boolean b
 - 9: let $b = (\forall i \in [0, n - 1] : X.P[i] \leq Y.P[i])$
 - 10: merge (X, Y) : payload Z
 - 11: let $\forall i \in [0, n - 1] : Z.P[i] = \max(X.P[i], Y.P[i])$
-

From *A comprehensive study of Convergent and Commutative Replicated Data Types* – Shapiro et al.
<https://hal.inria.fr/inria-00555588/document>



From *A comprehensive study of Convergent and Commutative Replicated Data Types* – Shapiro et al.
<https://hal.inria.fr/inria-00555588/document>

- **PNCounter** – Positive-Negative Counter
 - **Three operations:**
 - **Increment(x)**
 - **Decrement(x)**
 - **GetValue**

- **PNCounter** – Positive-Negative Counter
 - Cannot have global condition (e.g. $\text{PNCounter} > 0$)
 - Simple case – add 5 on one node, subtract (4 or 5) on another

- **PNCounter – Positive-Negative Counter**
 - **We can prevent decreasing if state on current node ≤ 0**
 - **But that means I cannot decrease even though I (globally) know that this should be possible**

- **PNCounter – Positive-Negative Counter**
- **One solution – system of „credits” / escrow transactions, etc.**

- **PNCounter – Positive-Negative Counter**
- **Another clever solution: When decrement is not possible because violating > 0 on local node, synchronise the cluster; but only then!**

- **ORSet – Observed-Remove Set**
 - **Set that allows for removal of elements that have been observed at the current node**
 - **Given causal support (version vectors) for merge operation and conflict resolution protocol that prefers additions over removals this is a CRDT**

- **ORSet** – Details beyond capacity of this presentation
 - **Read:** <https://arxiv.org/pdf/1210.3368v1.pdf>
An Optimized Conflict-free Replicated Set
Bieniusa, A. et al.

Specification 1 Outline of a state-based object specification. *Preconditions, arguments, return values and statements are optional.*

- 1: **payload** *Payload type; instantiated at all replicas*
 - 2: **initial** *Initial value*
 - 3: **query** *Query (arguments) : returns*
 - 4: **pre** *Precondition*
 - 5: **let** *Evaluate synchronously, no side effects*
 - 6: **update** *Source-local operation (arguments) : returns*
 - 7: **pre** *Precondition*
 - 8: **let** *Evaluate at source, synchronously*
 - 9: *Side-effects at source to execute synchronously*
 - 10: **compare** (value1, value2) : *boolean b*
 - 11: *Is $value1 \leq value2$ in semilattice?*
 - 12: **merge** (value1, value2) : *payload mergedValue*
 - 13: *LUB merge of value1 and value2, at any replica*
-

From *A comprehensive study of Convergent and Commutative Replicated Data Types* – Shapiro et al.
<https://hal.inria.fr/inria-00555588/document>

Specification 1 Outline of a state-based object specification. *Preconditions, arguments, return values and statements are optional.*

- 1: **payload** *Payload type; instantiated at all replicas*
- 2: **initial** *Initial value*
- 3: **query** *Query (arguments) : returns*
- 4: **pre** *Precondition*
- 5: **let** *Evaluate synchronously, no side effects*
- 6: **update** *Source-local operation (arguments)*
- 7: **pre** *Precondition*
- 8: **let** *Evaluate at source, no side effects*
- 9: *Side-effects at source, evaluate synchronously*
- 10: **compare** (value1, value2) *Compare two values*
- 11: *Is value1 ≤ value2 in the lattice?*
- 12: **merge** (value1, value2) *payload mergedValue*
- 13: *LUB merge of value1 and value2, at any replica*

Some 3 k LOC

Based on *A comprehensive study of Convergent and Commutative Replicated Data Types* – Shapiro et al.
<https://hal.inria.fr/inria-00555588/document>

- **You need top-level key to be able to access your objects:**

```
val DataKey = ORSetKey[String]("key")
```

- **Get**
- **To retrieve the current value of a data you send Replicator.Get message to the Replicator. You supply a consistency level which has the following meaning:**
 - **ReadLocal** the value will only be read from the local replica
 - **ReadFrom(n)** the value will be read and merged from n replicas, including the local replica
 - **ReadMajority** the value will be read and merged from a majority of replicas, i.e. at least $N/2 + 1$ replicas, where N is the number of nodes in the cluster (or cluster role group)
 - **ReadAll** the value will be read and merged from all nodes in the cluster (or all nodes in the cluster role group)

```
val replicator = DistributedData(system).replicator
val Counter1Key = PNCounterKey("counter1")
val Set1Key = GSetKey[String]("set1")
val Set2Key = ORSetKey[String]("set2")
```

```
replicator ! Get(Counter1Key, ReadLocal)
```

```
val readFrom3 = ReadFrom(n = 3, timeout = 1.second)
replicator ! Get(Set1Key, readFrom3)
```

```
val readMajority = ReadMajority(timeout = 5.seconds)
replicator ! Get(Set2Key, readMajority)
```



```
case g @ GetSuccess(Set1Key, req) =>
    val elements = g.get(Set1Key).elements
case GetFailure(Set1Key, req) =>
    // read from 3 nodes failed within 1.second
case NotFound(Set1Key, req) => // key set1 does not exist
```

- **Update**
- **To modify and replicate a data value you send a Replicator.Update message to the local Replicator.**
- You supply a write consistency level which has the following meaning:
 - **WriteLocal** the value will immediately only be written to the local replica, and later disseminated with gossip
 - **WriteTo(n)** the value will immediately be written to at least n replicas, including the local replica
 - **WriteMajority** the value will immediately be written to a majority of replicas, i.e. at least $N/2 + 1$ replicas, where N is the number of nodes in the cluster (or cluster role group)
 - **WriteAll** the value will immediately be written to all nodes in the cluster (or all nodes in the cluster role group)

```
val Counter1Key = PNCounterKey("counter1")
```

```
val Set1Key = GSetKey[String]("set1")
```

```
val Set2Key = ORSetKey[String]("set2")
```

```
replicator ! Update(Counter1Key, PNCounter(), WriteLocal)(_ + 1)
```

```
val writeTo3 = WriteTo(n = 3, timeout = 1.second)
```

```
replicator ! Update(Set1Key, GSet.empty[String], writeTo3)(_ + "hello")
```

```
val writeMajority = WriteMajority(timeout = 5.seconds)
```

```
replicator ! Update(Set2Key, ORSet.empty[String], writeMajority)(_ + "hello")
```

```
def receiveRemoveItem: Receive = {  
  case cmd @ RemoveItem(productId) =>  
    // Try to fetch latest from a majority of nodes first, since ORMap  
    // remove must have seen the item to be able to remove it.  
    replicator ! Get(DataKey, readMajority, Some(cmd))  
  
  case GetSuccess(DataKey, Some(RemoveItem(productId))) =>  
    replicator ! Update(DataKey, LWWMap(), writeMajority, None) {  
      _ - productId  
    }  
  case NotFound(DataKey, Some(RemoveItem(productId))) =>  
    // nothing to remove  
}
```

```
val s0 = GSet.empty[String]
val s1 = s0 + "a"
val s2 = s1 + "b" + "c"
if (s2.contains("a"))
    println(s2.elements) // a, b, c
```

```
implicit val node = Cluster(system)
val c0 = PNCounter.empty
val c1 = c0 + 1
val c2 = c1 + 7
val c3: PNCounter = c2 - 2
println(c3.value) // 6
```

Note – needs implicit node!

```
implicit val node = Cluster(system)
val s0 = ORSet.empty[String]
val s1 = s0 + "a"
val s2 = s1 + "b"
val s3 = s2 - "a"
println(s3.elements) // b
```

```
implicit val node = Cluster(system)
val m0 = ORMultiMap.empty[String, Int]
val m1 = m0 + ("a" -> Set(1, 2, 3))
val m2 = m1.addBinding("a", 4)
val m3 = m2.removeBinding("a", 2)
val m4 = m3.addBinding("b", 1)
println(m4.entries)
```


- Docs:
 - <http://doc.akka.io/docs/akka/2.5/scala/distributed-data.html>
- Interesting Specs from akka/akka
 - <https://github.com/akka/akka/tree/master/akka-distributed-data/src/multi-jvm/scala/akka/cluster/ddata>
- Interesting Implementations from akka/akka-samples (incl. ShoppingCart)
 - <https://github.com/akka/akka-samples/tree/master/akka-sample-distributed-data-scala/src/main/scala/sample/distributeddata>
- Interesting Specs from akka/akka-samples
 - <https://github.com/akka/akka-samples/tree/master/akka-sample-distributed-data-scala/src/multi-jvm/scala/sample/distributeddata>

Demo // FIXME

That's it...

GFT ■



Thank you! Questions? Answers?

https://en.wikipedia.org/wiki/Banner_Mania

私はカピバスも好きです