

SCALAWAVE.IO 2016

Applied Reactive Streaming Fast Data Processing with Akka

Jan Pustelnik @gosubpl 24 November 2016





Image source: Sinciair ZX Spectrum keys with BASIC keywords: Bill Bertram CC-B SA 2.5, via Wikimedia Commons.

GoSub strikes back ©









gosubpl / akka-online

Online and streaming algorithms with Akka

For good overview of streaming algos go to this debasishg's gist.

In org. apache. spank. streamDM project (the licencse is Apache 2.0 too). The code has been adapted to work with Akka streams instead of Spank streaming module by removing dependencies on the Spank stuff. Tested path is the HoeffdingTree model usage, which works best with input sourced from Arff files via SpecificationParser / ExampleParser - see see other parts of akka-online for usage examples. If you'd like to know more about HoeffdingTree on-line classifier, please read the HDT docs or go to Massive Online Analysis website that contains more information on the context.

Sample Arff files can be found in the MOA dataset repository.

For theoretical exposition to HoeffdingTree usage go to the original paper.

If you are asking yourself question why Akka, not Spark? - please read the classic bigger data, same laptop to see what we can gain going lightweight. In previous century, 4.5k records machine learning data set could be considered sizeable and hence all the growth of clustering for data processing. But with today's hardware I believe we can do better.

In the /lib directory of the project I have put suffixtree-1.0.0-SNAPSHOT.jar which is a compiled artifact of the Ukkonen's on-line Suffix Tree implementation. The licensse for this project is also Apache 2.0.

In this project I also use Guava implementation of Bloom Filter (see project dependencies) but there are other options.



- While preparing this slide deck I have received help from the wide Akka community:
 - all helpful people on akka/akka and akka/dev gitter chats
 - @ktoso, @hseeberger, @rbudzko Specifically

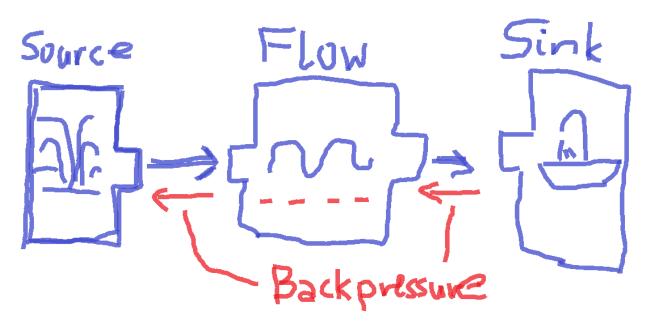












Backpressure ftw!





September 1981

RFC 792

Flow Control:

TCP provides a means for the receiver to govern the amount of data sent by the sender. This is achieved by returning a "window" with every ACK indicating a range of acceptable sequence numbers beyond the last segment successfully received. The window indicates an allowed number of octets that the sender may transmit before receiving further permission.

gateway may send a source quench message for message, the source host should cut back the rate at which it is sending traffic to the specified destination until it no longer receives source quench messages from the gateway. The source host can then gradually increase the rate at which it sends traffic to the destination until it again receives source quench messages.

The gateway or host may send the source quench message when it approaches its capacity limit rather than waiting until the capacity is exceeded. This means that the data datagram which triggered the source quench message may be delivered.

Staged event-driven architecture, ftw! (2001)





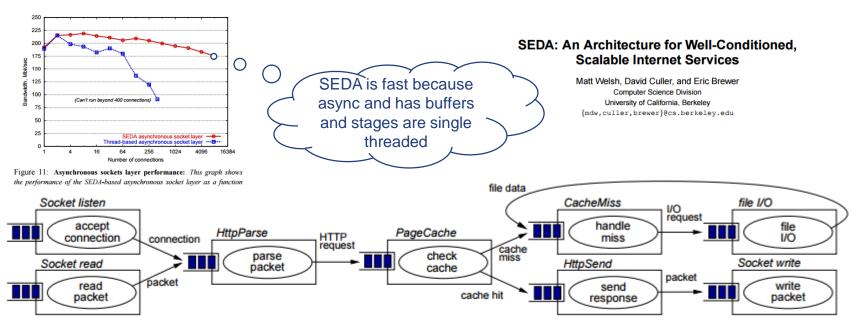


Figure 5: **Staged event-driven (SEDA) HTTP server:** This is a structural representation of the SEDA-based Web server, described in detail in Section 5.1. The application is composed as a set of stages separated by queues. Edges represent the flow of events between stages. Each stage can be independently managed, and stages can be run in sequence or in parallel, or a combination of the two. The use of event queues allows each stage to be individually load-conditioned, for example, by thresholding its event queue. For simplicity, some event paths and stages have been elided from this figure.



On-Line Algorithms Versus Off-Line Algorithms: How Much is it Worth to Know the Future?

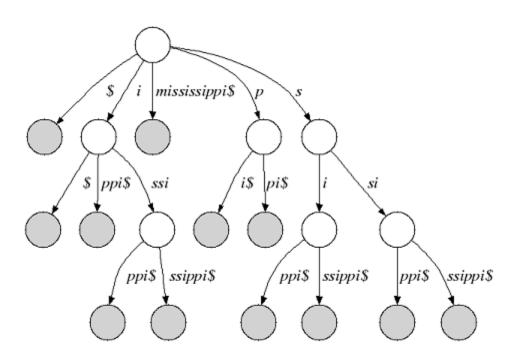


Richard M. Karp* TR-92-044 July 1992

Abstract

An on-line algorithm is one that receives a sequence of requests and performs an immediate action in response to each request. On-line algorithms arise in any situation where decisions must be made and resources allocated without knowledge of the future. The effectiveness of an on-line algorithm may be measured by its competitive ratio, defined as the worst-case ratio between its cost and that of a hypothetical off-line algorithm which knows the entire sequence of requests in advance and chooses its actions optimally. In a variety of settings, we discuss techniques for proving upper and lower bounds on the competitive ratios achievable by on-line algorithms. In particular, we discuss the advantages of randomized on-line algorithms over deterministic ones.





Source: http://docs.seqan.de/seqan/1.2/PAGE_Indices.html

Ukkonen algorithm for on-line Suffix tree construction





Algorithmica (1995) 14: 249-260

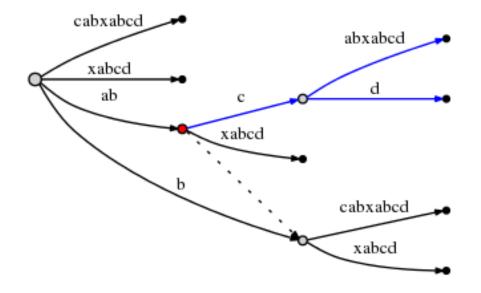


On-Line Construction of Suffix Trees1

E. Ukkonen²

Abstract. An on-line algorithm is presented for constructing the suffix tree for a given string in time linear in the length of the string. The new algorithm has the desirable property of processing the string symbol by symbol from left to right. It always has the suffix tree for the scanned part of the string ready. The method is developed as a linear-time version of a very simple algorithm for (quadratic size) suffix tries. Regardless of its quadratic worst case this latter algorithm can be a good practical method when the string is not too long. Another variation of this method is shown to give, in a natural way, the well-known algorithms for constructing suffix automata (DAWGs).

Key Words. Linear-time algorithm, Suffix tree, Suffix trie, Suffix automaton, DAWG.



Source: http://stackoverflow.com/questions/9452701/ukkonens-suffix-tree-algorithm-in-plain-english



 Build on-line generalized suffix tree based on the genome signatures of known contaminants

{aggtca, attgca, gatttaca}

 Look-up whether our specimen does not contain a substring common with one of the known contaminants

```
ttg -> {aggtca, attgca, gatttaca}
```

- There can be new known contaminants... add them as they go
- Repeat search if necessary...







Streaming algorithms





Jérémie Lumbroso's talk at the AK Data Science Summit on Streaming and Sketching in Big Data and Analytics.

Source: https://speakerdeck.com/timonk/philippe-flajolets-contribution-to-streaming-algorithms

- ONE PASS (on-line algorithm)
- constant/logarithmic memory

...

■ 1970, 1983 ... old stuff ©

For more: https://gist.github.com/debasishg/8172796

Historical context

- ▶ 1970: average-case \rightarrow deterministic algorithms on random input
- ▶ 1976-78: first randomized algorithms (primality testing, matrix multiplication verification, find nearest neighbors)
- ▶ **1979**: Munro and Paterson, find median in one pass with $\Theta(\sqrt{n})$ space with high probability
 - ⇒ (almost) first streaming algorithm

In **1983**, <u>Probabilistic Counting</u> by Flajolet and Martin is (more or less) the first streaming algorithm (one pass + constant/logarithmic memory).

Google scholar

Probabilistic counting algorithms for data base applications

P Flajelst...-Journal of computer and system sciences, 1985. Elizavier Abstract This paper introduces a class of probabilistic conting algorithms with which one can estimate the number of distinct elements in a large collection of data dypically a large file stored on disk) in a sngle pass using only a small additional storage (typically less ... Circle by 628 - Binited articles - A 38 x zerooms

Probabilistic counting

P Flajolet. - Foundstons of Computer Science. ..., 1983 - isoexploresices org Abstract We present here a class of probabilistic algorithms with which one can estimate the number of distinct elements in a collection of data (typically a large file stored on disk.) in a single pass, using only 0 (1) auxiliary storage and 0 (1) operations per element. We ... Clote by 111 - Related articles - All Yearsions

Combining both versions: cited about 750 times = **second most cited** element of Philippe's bibliography, after only *Analytic Combinatorics*.

4/22

Example: Kadane algorithm (on-line, streaming, 1984)



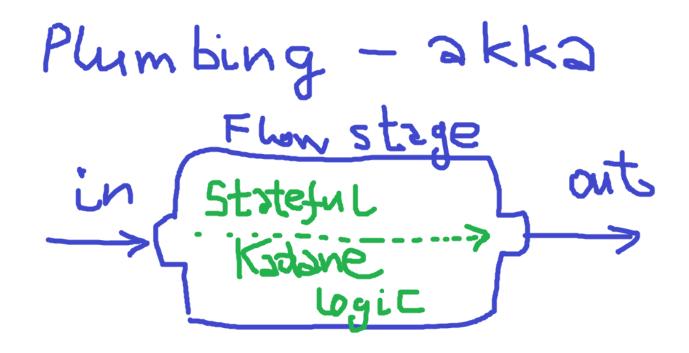
(...) the **maximum subarray problem** is the task of finding the contiguous subarray within a one-dimensional array of numbers which has the largest sum. For example, for the sequence of values -2, 1, -3, 4, -1, 2, 1, -5, 4; the contiguous subarray with the largest sum is 4, -1, 2, 1, with sum 6.

```
def max_subarray(A):
    max_ending_here = max_so_far = 0
    for x in A:
        max_ending_here = max(0, max_ending_here + x)
        max_so_far = max(max_so_far, max_ending_here)
    return max_so_far
```

Source: https://en.wikipedia.org/wiki/Maximum_subarray_problem









```
// Shape definition
val in: Inlet[Int] = Inlet("KadaneFlowStage.in")
val out: Outlet[Int] = Outlet("KadaneFlowStage.out")
override val shape: FlowShape[Int, Int] = FlowShape(in, out)
            ---> >Inlet > Logic > Outlet> --->
```



```
// Handler for the Outlet
setHandler(out, new OutHandler {
  override def onPull(): Unit = {
    if (!hasBeenPulled(in))
      pull(in)
  }
})
```



```
// Handler for the Outlet
setHandler(out, new OutHandler {
 override def onPull(): Unit = {
 PDD - Pull Driven Development
```

Flow Shape – output handler (plumbing)



```
// Handler for the Outlet
setHandler(out, new OutHandler {
```

Flow stage - "Business logic"



```
new GraphStageLogic(shape) {
 // state
 var maxEndingHere = 0
 var maxSoFar = 0
 // Handler(s) for the Inlet
  setHandler(in, new InHandler {
   // what to do when a new element is ready to be consumed
   override def onPush(): Unit = {
     val elem = grab(in)
     // "Business" logic
     maxEndingHere = Math.max(0, maxEndingHere + elem)
     maxSoFar = Math.max(maxSoFar, maxEndingHere)
     // this should never happen
     // we decide to not push the value, avoiding the error
     // but potentially losing the value
     if (isAvailable(out))
        push(out, maxSoFar)
```



Flow stage - "Business logic"



```
new GraphStageLogic(shape) {
  // state
 var maxEndingHere = 0
 var maxSoFar = 0
 // Handler(s) for the Inlet
  setHandler(in, new InHandler {
   // what to do when a new element is ready to be consumed
   override def onPush(): Unit = {
     val elem = grab(in)
     // "Business" logic
     maxEndingHere = Math.max(0, maxEndingHere + elem)
     maxSoFar = Math.max(maxSoFar, maxEndingHere)
     // this should never happen
     // we decide to not push the value, avoiding the error
     // but potentially losing the value
     if (isAvailable(out))
        push(out, maxSoFar)
```

Proper Kadane algo logic here



Let the flow flow...

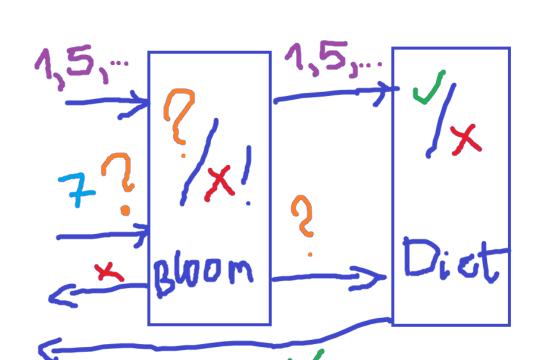


```
object KadaneFlowMain extends App {
 implicit val system = ActorSystem()
 implicit val mat = ActorMaterializer()
 val kadaneFlowStage = new KadaneFlowStage
 val done = Source.repeat(1).take(100).map(_ => Random.nextInt(1100) - 1000)
     .throttle(1, Duration(100, "millisecond"), 1, ThrottleMode.shaping)
    .via(kadaneFlowStage)
    .throttle(1, Duration(100, "millisecond"), 1, IhrottleMode.shaping)
    .runWith(Sink.foreach(println))
 import scala.concurrent.ExecutionContext.Implicits.global
 done.onComplete( => system.terminate())
 Await.ready(system.whenTerminated, Duration.Inf)
```



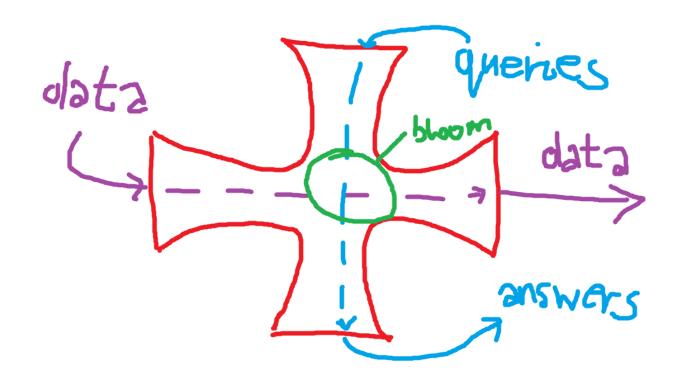




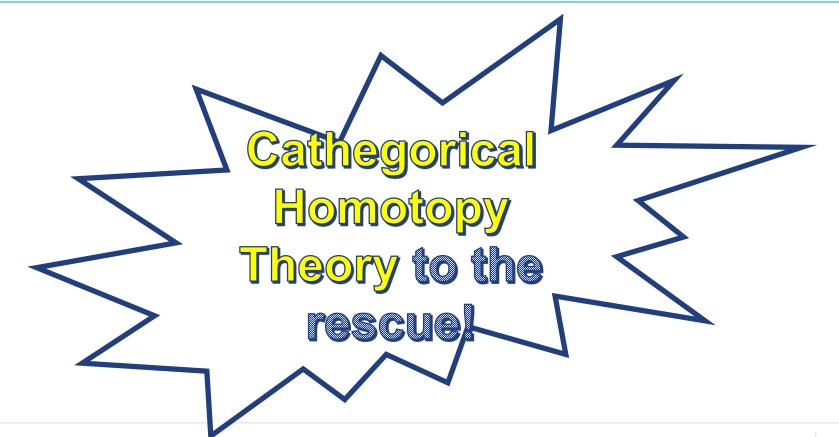


A good shape for bloom filter? CrossShape!









Remember, a shape is just a shape

out2: Outlet[Out2 @uncheckedVariance]) extends Shape {



```
/**
 * A bidirectional flow of elements that consequently has two inputs and two
 * outputs, arranged like this:
 *
                                                         In2
 * {{{
 * In1 ~> | ~> Out1
                                             In1 ~>
                                                        cross
                                                                 ~> Out1
         | bidi |
 * Out2 <~ | <~ In2
                                                         Out2
 * }}}
final case class BidiShape[-In1, +Out1, -In2, +Out2](
  in1: Inlet[In1 @uncheckedVariance],
  out1: Outlet[Out1 @uncheckedVariance],
  in2: Inlet[In2 @uncheckedVariance],
```

BloomFilterCrossStage, ftw!



```
// Cross Shape is actually BidiShape - for shape the semantics doesn't count, only syntax
class BloomFilterCrossStage extends GraphStage[BidiShape[Int, Int, Int, String]] {
    // Stage syntax
    val dataIn: Inlet[Int] = Inlet("BloomFilterCrossStage.dataIn")
    val dataOut: Outlet[Int] = Outlet("BloomFilterCrossStage.dataOut")
    val queriesIn: Inlet[Int] = Inlet("BloomFilterCrossStage.queriesIn")
    val answersOut: Outlet[String] = Outlet("BloomFilterCrossStage.answersOut")
    override val shape: BidiShape[Int, Int, Int, String] = BidiShape(dataIn, dataOut, queriesIn, answersOut)
```

Two crossing flows...

```
override def onPush(): Unit = {
                                            val x = grab(queriesIn)
                                            val answer = if (filter.mightContain(x)) {
                                             s"MAYBE, filter probably contains $x"
                                            } else {
                                             s"NO, filter definitely does not contain $x"
                                            if (isAvailable(answersOut))
setHandler(dataIn, new InHandler {
                                              push(answersOut, answer)
  override def onPush(): Unit = {
    val elem = grab(dataIn)
    filter.put(elem)
    if (isAvailable(dataOut))
      push(dataOut, elem)
                               setHandler(answersOut, new OutHandler {
                                 override def onPull(): Unit = {
                                    if (!hasBeenPulled(queriesIn))
                                       pull(queriesIn)
```

setHandler(queriesIn, new InHandler {

```
setHandler(dataOut, new OutHandler {
  override def onPull(): Unit = {
    if (!hasBeenPulled(dataIn))
      pull(dataIn)
  }
})
```



Akka allows you to put your **on-line / streaming data structure / algorithm** in context. You don't need to think about how to take care of data flow and backpressure.

Well thought out architecture of **Akka** lets you **concentrate on the stuff** *relevant to your problem domain*.

Akka is geared towards high performance, and can be used in IoT setup, where e.g. Spark streaming would not fit.

It is **easy to port** e.g. machine learning **algorithms from other streaming setups** (like Spark streaming) **to Akka**.

ProTip #1 – you can use actor to better manage state for on-line algos



```
val actorQueue = Source.queue[Int](10, OverflowStrategy.backpressure)
val actorSink = actorQueue
  .throttle(1, Duration(100, "millisecond"), 1, ThrottleMode.shaping)
  .to(Sink.foreach((x) => println(x))).run()
val done = actorSink.watchCompletion()
val kadaneFlowActor = system.actorOf(Props(new KadaneFlowActor(actorSink)))
Source.repeat(1).take(100).map( => Random.nextInt(1100) - 1000)
  .runWith(Sink.actorRefWithAck(kadaneFlowActor, "TEST", "ACK", "END", => "FAIL"))
```





```
case elem : Int => // onPush() + grab(in)
 // "Business" logic
 max ending here = Math.max(0, max ending here + elem)
 max so far = Math.max(max so far, max ending here)
   Thread.sleep(100) // FIXME: don't do this at home
 log.info(s"Received: $elem, sending out: $max so far")
 val offF = queue.offer(max so far) // push element downstream
                                 // generate backpressure in the Actor
 offF pipeTo self
case e : QueueOfferResult =>
 upStream ! "ACK" // ask for next element
```

ProTip #2 – it is (too) easy to create new shapes



```
/**
 * A Y-shaped flow of elements that consequently has two inputs
 * and one output, arranged like this:
 *
 * {{{
 * +----+
 * In1 ~>
   | tripod |~> Out
  In2 ~>|
  +----+
 * }}}
```

Tripod? Just like the in old days ©



```
final case class TripodShape[-In1, -In2, +Out](
 in1: Inlet[In1 @uncheckedVariance],
 in2: Inlet[In2 @uncheckedVariance],
 out: Outlet[Out @uncheckedVariance]) extends Shape {
   override val inlets: immutable.Seq[Inlet[ ]] = List(in1, in2)
   override val outlets: immutable.Seq[Outlet[_]] = List(out)
   override def deepCopy(): TripodShape[In1, In2, Out] =
     TripodShape(in1.carbonCopy(), in2.carbonCopy(), out.carbonCopy())
   override def copyFromPorts(inlets: immutable.Seq[Inlet[]], outlets: immutable.Seq[Outlet[]]): Shape = {
      require(inlets.size == 2, s"proposed inlets [${inlets.mkString(", ")}] do not fit TripodShape")
      require(outlets.size == 1, s"proposed outlets [${outlets.mkString(", ")}] do not fit TripodShape")
     TripodShape(inlets(0), inlets(1), outlets(0))
 def reversed: Shape = copyFromPorts(inlets.reverse, outlets.reverse)
```

Remember your Topology class? A shape is just a shape...





ProTip #3 – you can return materialized value from GraphDSL built shape



```
val crossStage = new BloomFilterCrossStage
val graph = RunnableGraph.fromGraph(GraphDSL.create(Sink.foreach(println)) { implicit builder => outControl =>
  import GraphDSL.Implicits.
  val inData = Source.repeat(1).take(100).map( => Random.nextInt(1000)).throttle(1, Duration(100, "millisecond")
  val outData = Sink.foreach(println)
  val inControl = Source.repeat(1).take(10).map( => Random.nextInt(2000) - 1000).throttle(1, Duration(1500, "m
  //val outControl = Sink.foreach(println) // Moved to foreach/builder construct
  val cross = builder.add(crossStage)
  inData  cross.in1; cross.out1  outData
  inControl ~> cross.in2; cross.out2 ~> outControl
  ClosedShape
}).run()
```









Chank You! Questions? Answers

https://en.wikipedia.org/wiki/Banner_Mania